

Adam Harries

CUDA, or how I stopped worrying and learned to love the GPU

Tuesday October 15th 2013

Overview

The talk

The talk

What this talk will be:

The talk

What this talk will be:

- ▶ Some history (of both GPUs and CUDA)

The talk

What this talk will be:

- ▶ Some history (of both GPUs and CUDA)
- ▶ A basic introduction to CUDA, how it works, and using it in Python

The talk

What this talk will be:

- ▶ Some history (of both GPUs and CUDA)
- ▶ A basic introduction to CUDA, how it works, and using it in Python
- ▶ Some basic examples of speeding up existing Python programs using CUDA

The talk

What this talk will be:

- ▶ Some history (of both GPUs and CUDA)
- ▶ A basic introduction to CUDA, how it works, and using it in Python
- ▶ Some basic examples of speeding up existing Python programs using CUDA

The talk

What this talk will be:

- ▶ Some history (of both GPUs and CUDA)
- ▶ A basic introduction to CUDA, how it works, and using it in Python
- ▶ Some basic examples of speeding up existing Python programs using CUDA

What this talk won't be

The talk

What this talk will be:

- ▶ Some history (of both GPUs and CUDA)
- ▶ A basic introduction to CUDA, how it works, and using it in Python
- ▶ Some basic examples of speeding up existing Python programs using CUDA

What this talk won't be

- ▶ A deep, massively complicated exploration of the CUDA programming model

The talk

What this talk will be:

- ▶ Some history (of both GPUs and CUDA)
- ▶ A basic introduction to CUDA, how it works, and using it in Python
- ▶ Some basic examples of speeding up existing Python programs using CUDA

What this talk won't be

- ▶ A deep, massively complicated exploration of the CUDA programming model
- ▶ An in-depth guide to optimizing your CUDA programs (see Aidan Chalk about that)

The talk

What this talk will be:

- ▶ Some history (of both GPUs and CUDA)
- ▶ A basic introduction to CUDA, how it works, and using it in Python
- ▶ Some basic examples of speeding up existing Python programs using CUDA

What this talk won't be

- ▶ A deep, massively complicated exploration of the CUDA programming model
- ▶ An in-depth guide to optimizing your CUDA programs (see Aidan Chalk about that)
- ▶ A lecture on parallel algorithms

The talk

What this talk will be:

- ▶ Some history (of both GPUs and CUDA)
- ▶ A basic introduction to CUDA, how it works, and using it in Python
- ▶ Some basic examples of speeding up existing Python programs using CUDA

What this talk won't be

- ▶ A deep, massively complicated exploration of the CUDA programming model
- ▶ An in-depth guide to optimizing your CUDA programs (see Aidan Chalk about that)
- ▶ A lecture on parallel algorithms
- ▶ Coherent

History

First, some history

GPU - Speeding up graphics computation since 1983

First, some history

GPU - Speeding up graphics computation since 1983

- ▶ Dedicated cards for handling graphical rendering

First, some history

GPU - Speeding up graphics computation since 1983

- ▶ Dedicated cards for handling graphical rendering
- ▶ Generally 3d acceleration today, but started off with 2d acceleration

First, some history

GPU - Speeding up graphics computation since 1983

- ▶ Dedicated cards for handling graphical rendering
- ▶ Generally 3d acceleration today, but started off with 2d acceleration
- ▶ Arguably first card: Intel iSBX 275 Video Graphics Controller Multimodule Board

Intel iSBX 275

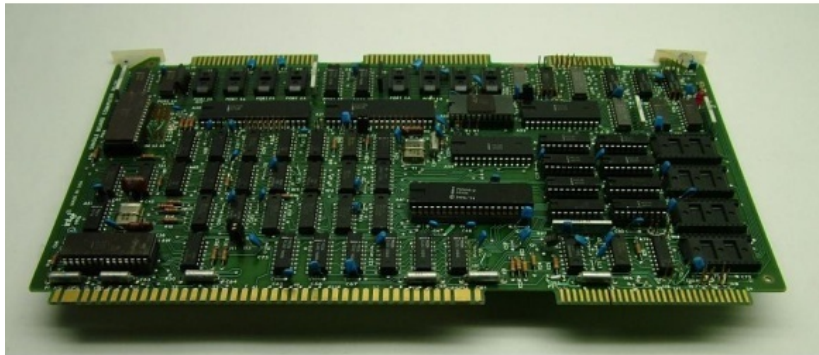


Figure : The Intel ISBX, because everyone loves dual in line packaging...

The arrival of 3d

Dedicated 3d graphics hardware speedup

The arrival of 3d

Dedicated 3d graphics hardware speedup

- ▶ 3dfx interactive's "Voodoo 1" card

The arrival of 3d

Dedicated 3d graphics hardware speedup

- ▶ 3dfx interactive's "Voodoo 1" card
- ▶ Dedicated 3d graphics card

The arrival of 3d

Dedicated 3d graphics hardware speedup

- ▶ 3dfx interactive's "Voodoo 1" card
- ▶ Dedicated 3d graphics card
- ▶ needed to be piggybacked with a 2d card

Voodoo 1 card

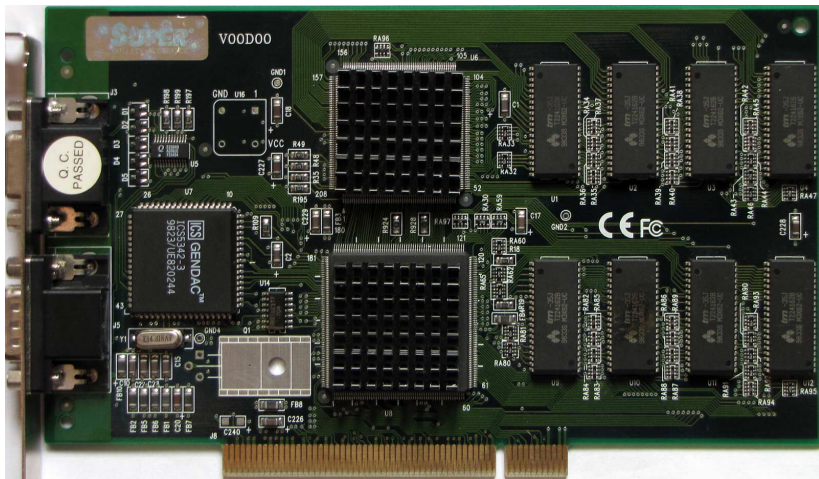


Figure : The 3dfx voodoo 1 card. Note the dual VGA ports, for piggybacking off a 2d card.

Programmable pipeline

The advent of programming on the graphics card

- ▶ Programmers could now define “shader programs”

Programmable pipeline

The advent of programming on the graphics card

- ▶ Programmers could now define “shader programs”
- ▶ First graphics card with support: nVidia GeForce 3

Programmable pipeline

The advent of programming on the graphics card

- ▶ Programmers could now define “shader programs”
- ▶ First graphics card with support: nVidia GeForce 3
- ▶ Only 12 years ago - in 2001

nVidia Geforce 3

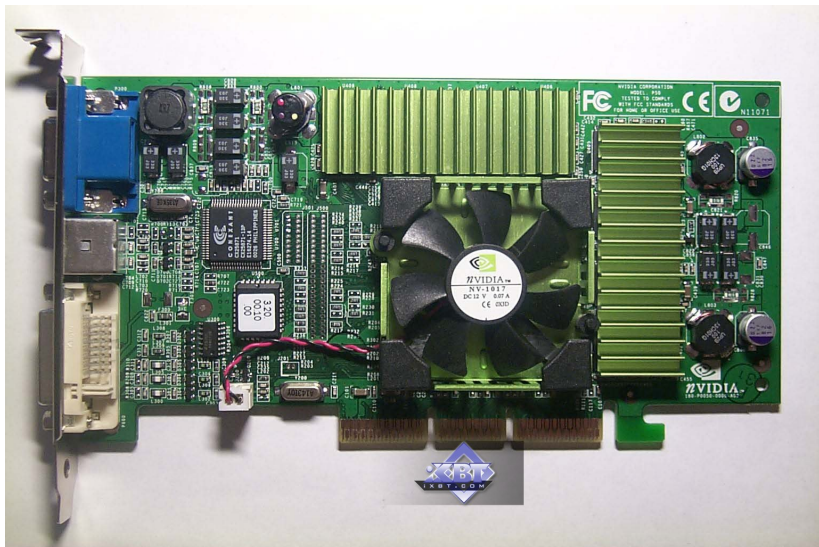


Figure : The first graphics card with a programmable pipeline

CUDA arrives!

- ▶ Only available on nVidia cards (obviously)

CUDA arrives!

- ▶ Only available on nVidia cards (obviously)
- ▶ Released to public in

CUDA arrives!

- ▶ Only available on nVidia cards (obviously)
- ▶ Released to public in
- ▶ First supporting card: 8800GTX

8800GTX



CUDA: an overview

Kernels

- ▶ CUDA is based upon the concept of kernels

Kernels

- ▶ CUDA is based upon the concept of kernels
- ▶ similar to functions

Kernels

- ▶ CUDA is based upon the concept of kernels
- ▶ similar to functions
- ▶ written in a dialect of C++

Kernels

- ▶ CUDA is based upon the concept of kernels
- ▶ similar to functions
- ▶ written in a dialect of C++
- ▶ run on the GPU

Kernels

- ▶ CUDA is based upon the concept of kernels
- ▶ similar to functions
- ▶ written in a dialect of C++
- ▶ run on the GPU
- ▶ one kernel per thread

Kernels

- ▶ CUDA is based upon the concept of kernels
- ▶ similar to functions
- ▶ written in a dialect of C++
- ▶ run on the GPU
- ▶ one kernel per thread

Kernels

- ▶ CUDA is based upon the concept of kernels
- ▶ similar to functions
- ▶ written in a dialect of C++
- ▶ run on the GPU
- ▶ one kernel per thread

Example:

```
1  __global__ void double(int *a)
2  {
3      const int id = threadIdx.x;
4      a[id] = 2*a[id];
5  }
```


Supplementary functions

- ▶ sometimes we can't fit everything in a single function

Supplementary functions

- ▶ sometimes we can't fit everything in a single function
- ▶ use `__device__` functions, callable from other function running on GPU

Supplementary functions

- ▶ sometimes we can't fit everything in a single function
- ▶ use `__device__` functions, callable from other function running on GPU
- ▶ `__global__` functions can't be called from anything, apart from invocation

Supplementary functions

- ▶ sometimes we can't fit everything in a single function
- ▶ use `__device__` functions, callable from other function running on GPU
- ▶ `__global__` functions can't be called from anything, apart from invocation
- ▶ `__device__` functions can't be invoked by the CPU, must be called by another function on GPU

Supplementary functions

- ▶ sometimes we can't fit everything in a single function
- ▶ use `__device__` functions, callable from other function running on GPU
- ▶ `__global__` functions can't be called from anything, apart from invocation
- ▶ `__device__` functions can't be invoked by the CPU, must be called by another function on GPU

Supplementary functions

- ▶ sometimes we can't fit everything in a single function
- ▶ use `__device__` functions, callable from other function running on GPU
- ▶ `__global__` functions can't be called from anything, apart from invocation
- ▶ `__device__` functions can't be invoked by the CPU, must be called by another function on GPU

```
1  __device__ int square(int x)
2  {
3      return x*x;
4  }
5  __global__ void double_and_square(int *a)
6  {
7      const int id = threadIdx.x;
8      a[id] = 2*square(a[id]);
9  }
```

Running lots of kernels at once

- ▶ CUDA uses “thread blocks” to achieve parallelism

Running lots of kernels at once

- ▶ CUDA uses “thread blocks” to achieve parallelism
- ▶ define a dimension for the blocks running on

Running lots of kernels at once

- ▶ CUDA uses “thread blocks” to achieve parallelism
- ▶ define a dimension for the blocks running on
- ▶ define a dimension for threads within blocks

Running lots of kernels at once

- ▶ CUDA uses “thread blocks” to achieve parallelism
- ▶ define a dimension for the blocks running on
- ▶ define a dimension for threads within blocks
- ▶ limitations depending on card being used

Running lots of kernels at once

- ▶ CUDA uses “thread blocks” to achieve parallelism
- ▶ define a dimension for the blocks running on
- ▶ define a dimension for threads within blocks
- ▶ limitations depending on card being used
- ▶ use variables based on dimensions to tell kernel what ID it is

Some examples

An old lab

Visualising the interference of waves using a discrete grid of cells

An old lab

Visualising the interference of waves using a discrete grid of cells

- ▶ this is an actual exercise from last years Year 1 physics lab module

An old lab

Visualising the interference of waves using a discrete grid of cells

- ▶ this is an actual exercise from last years Year 1 physics lab module
- ▶ I haven't done it how it was recommended by the physics department

An old lab

Visualising the interference of waves using a discrete grid of cells

- ▶ this is an actual exercise from last years Year 1 physics lab module
- ▶ I haven't done it how it was recommended by the physics department
- ▶ nested for Loops over an array to calculate the value at each cell

An old lab

Visualising the interference of waves using a discrete grid of cells

- ▶ this is an actual exercise from last years Year 1 physics lab module
- ▶ I haven't done it how it was recommended by the physics department
- ▶ nested for Loops over an array to calculate the value at each cell
- ▶ both versions (cuda and traditional) use the same formula to calculate the value at each cell, from each wave (sin of hypotenuse):

An old lab

Visualising the interference of waves using a discrete grid of cells

- ▶ this is an actual exercise from last years Year 1 physics lab module
- ▶ I haven't done it how it was recommended by the physics department
- ▶ nested for Loops over an array to calculate the value at each cell
- ▶ both versions (cuda and traditional) use the same formula to calculate the value at each cell, from each wave (sin of hypotenuse):

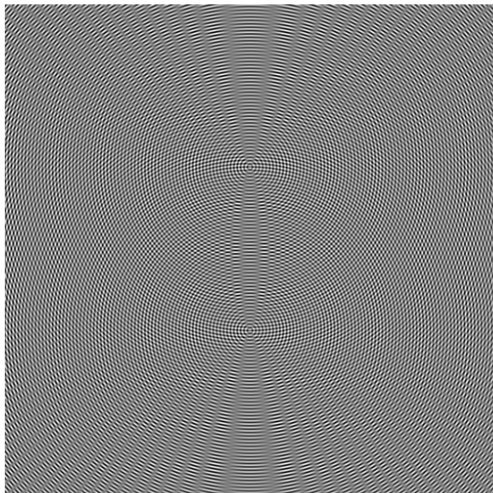
An old lab

Visualising the interference of waves using a discrete grid of cells

- ▶ this is an actual exercise from last years Year 1 physics lab module
- ▶ I haven't done it how it was recommended by the physics department
- ▶ nested for Loops over an array to calculate the value at each cell
- ▶ both versions (cuda and traditional) use the same formula to calculate the value at each cell, from each wave (sin of hypotenuse):

$$Wavevalue = \sin(\sqrt{(x - s_x)^2 + (y - s_y)^2})$$

The result



“Tradtional” loops version:

```
1  from numpy import *
2  import matplotlib.pyplot as plt
3  wav_array = zeros((300,300), dtype=float32)
4  pa = (300,450)
5  pb = (600,450)
6  for i in range(300):
7      for j in range(300):
8          adist = sin( sqrt( (i-pa[0])**2 + (j-pa[1])**2 ))
9          bdist = sin( sqrt( (i-pb[0])**2 + (j-pb[1])**2 ))
10         wav_array[i][j] = (adist+bdist)/2
11  plt.imshow(wav_array)
12  plt.clim(-1,1);
13  plt.set_cmap('gray')
14  plt.axis('off')
15  plt.show()
```

And some analysis

- ▶ not that fast - $O(n^2)$ in fact

And some analysis

- ▶ not that fast - $O(n^2)$ in fact
- ▶ still not as slow as it could be

And some analysis

- ▶ not that fast - $O(n^2)$ in fact
- ▶ still not as slow as it could be
- ▶ still not as fast either (eg, could be in C)

And some analysis

- ▶ not that fast - $O(n^2)$ in fact
- ▶ still not as slow as it could be
- ▶ still not as fast either (eg, could be in C)
- ▶ timings (from unix time utility, run on my laptop):

And some analysis

- ▶ not that fast - $O(n^2)$ in fact
- ▶ still not as slow as it could be
- ▶ still not as fast either (eg, could be in C)
- ▶ timings (from unix time utility, run on my laptop):

And some analysis

- ▶ not that fast - $O(n^2)$ in fact
- ▶ still not as slow as it could be
- ▶ still not as fast either (eg, could be in C)
- ▶ timings (from unix time utility, run on my laptop):

real 0m12.579s

user 0m10.220s

sys 0m0.070s

- ▶ still faster than the way recommended by the physics department

The CUDA version

- ▶ the CUDA version is more difficult to display in a talk

The CUDA version

- ▶ the CUDA version is more difficult to display in a talk
- ▶ a little longer - but due to structuring, not difficulty

The CUDA version

- ▶ the CUDA version is more difficult to display in a talk
- ▶ a little longer - but due to structuring, not difficulty
- ▶ comprised of two main parts: kernel code, and calling/structure (CPU) code

The CUDA version

- ▶ the CUDA version is more difficult to display in a talk
- ▶ a little longer - but due to structuring, not difficulty
- ▶ comprised of two main parts: kernel code, and calling/structure (CPU) code
- ▶ all CPU code in Python

The CUDA version

- ▶ the CUDA version is more difficult to display in a talk
- ▶ a little longer - but due to structuring, not difficulty
- ▶ comprised of two main parts: kernel code, and calling/structure (CPU) code
- ▶ all CPU code in Python
- ▶ kernel written in CUDA (somewhat like C)

The structure/CPU code

```
1  from numpy import *
2  import matplotlib.pyplot as plt
3  import pycuda.compiler as comp
4  import pycuda.driver as drv
5  import pycuda.autoinit
6  mod = comp.SourceModule("""
7  #kernel declaration, of "wave_kernel"
8  """)
9  wav_array = zeros((900,900), dtype=float32)
10 wav_array_gpu = drv.mem_alloc(wav_array.nbytes)
11 make_waves = mod.get_function("wave_kernel")
12 make_waves(wav_array_gpu,block=(30,30,1), grid=(30,30))
13 drv.memcpy_dtoh(wav_array, wav_array_gpu)
```

Walking through the code

Generic imports, with the most interesting ones:

Walking through the code

Generic imports, with the most interesting ones:

```
1 import pycuda.compiler as comp
2 import pycuda.driver as drv
3 import pycuda.autoinit
```

Walking through the code

Generic imports, with the most interesting ones:

```
1 import pycuda.compiler as comp
2 import pycuda.driver as drv
3 import pycuda.autoinit
```

They respectively:

Walking through the code

Generic imports, with the most interesting ones:

```
1 import pycuda.compiler as comp
2 import pycuda.driver as drv
3 import pycuda.autoinit
```

They respectively:

- ▶ give access to nvcc (nvidia cuda compiler)

Walking through the code

Generic imports, with the most interesting ones:

```
1 import pycuda.compiler as comp
2 import pycuda.driver as drv
3 import pycuda.autoinit
```

They respectively:

- ▶ give access to nvcc (nvidia cuda compiler)
- ▶ interface with CUDA, and provide a host of useful abstractions

Walking through the code

Generic imports, with the most interesting ones:

```
1 import pycuda.compiler as comp
2 import pycuda.driver as drv
3 import pycuda.autoinit
```

They respectively:

- ▶ give access to nvcc (nvidia cuda compiler)
- ▶ interface with CUDA, and provide a host of useful abstractions
- ▶ initialise much of what CUDA does, including getting kernels ready to launch

More walking

```
1 mod = comp.SourceModule("""  
2 #kernel declaration, of "wave_kernel"  
3 """)
```


More walking

```
1  mod = comp.SourceModule("""
2  #kernel declaration, of "wave_kernel"
3  """)

4  wav_array = zeros((900,900), dtype=float32)
5  wav_array_gpu = drv.mem_alloc(wav_array.nbytes)
```

More walking

```
1  mod = comp.SourceModule("""
2  #kernel declaration, of "wave_kernel"
3  """)

4  wav_array = zeros((900,900), dtype=float32)
5  wav_array_gpu = drv.mem_alloc(wav_array.nbytes)

6  make_waves = mod.get_function("wave_kernel")
7  make_waves(wav_array_gpu,block=(30,30,1), grid=(30,30))
```

More walking

```
1  mod = comp.SourceModule("""
2  #kernel declaration, of "wave_kernel"
3  """)

4  wav_array = zeros((900,900), dtype=float32)
5  wav_array_gpu = drv.mem_alloc(wav_array.nbytes)

6  make_waves = mod.get_function("wave_kernel")
7  make_waves(wav_array_gpu,block=(30,30,1), grid=(30,30))

8  drv.memcpy_dtoh(wav_array, wav_array_gpu)
```

The kernel code

```
1  __global__ void wave_kernel(float *wa)
2  {
3      const int j = threadIdx.y+(blockIdx.y*gridDim.y);
4      const int i = threadIdx.x+(blockIdx.x*gridDim.x);
5      const int lID = i+(j*blockDim.y*gridDim.y);
6      float pa[] = {300,450};
7      float pb[] = {600,450};
8      float a,b;
9      a = sqrt(((i-pa[0])*(i-pa[0]))+((j-pa[1])*(j-pa[1])));
10     b = sqrt(((i-pb[0])*(i-pb[0]))+((j-pb[1])*(j-pb[1])));
11     wa[lID] = (sin(a)+sin(b))/2;
12 }
```

Strolling through the kernel

```
1  const int j = threadIdx.y+(blockIdx.y*gridDim.y);  
2  const int i = threadIdx.x+(blockIdx.x*gridDim.x);  
3  const int lID = i+(j*blockDim.y*gridDim.y);
```

Strolling through the kernel

```
1  const int j = threadIdx.y+(blockIdx.y*gridDim.y);  
2  const int i = threadIdx.x+(blockIdx.x*gridDim.x);  
3  const int lID = i+(j*blockDim.y*gridDim.y);
```

- ▶ Threads getting information about themselves

Strolling through the kernel

```
1  const int j = threadIdx.y+(blockIdx.y*gridDim.y);  
2  const int i = threadIdx.x+(blockIdx.x*gridDim.x);  
3  const int lID = i+(j*blockDim.y*gridDim.y);
```

- ▶ Threads getting information about themselves
- ▶ Allows threads to determine which parts of problems they should work on

Strolling through the kernel

```
1  const int j = threadIdx.y+(blockIdx.y*gridDim.y);  
2  const int i = threadIdx.x+(blockIdx.x*gridDim.x);  
3  const int lID = i+(j*blockDim.y*gridDim.y);
```

- ▶ Threads getting information about themselves
- ▶ Allows threads to determine which parts of problems they should work on
- ▶ in this case, (i,j) is the cell, while lID is the element of the final array

Some analysis

- ▶ Arguably an order of magnitude faster than the serial method

Some analysis

- ▶ Arguably an order of magnitude faster than the serial method
- ▶ Still some slowdowns in copying to/from GPU, but shouldn't be a problem if clever about it, and not constantly transferring data

Some analysis

- ▶ Arguably an order of magnitude faster than the serial method
- ▶ Still some slowdowns in copying to/from GPU, but shouldn't be a problem if clever about it, and not constantly transferring data
- ▶ Slow to invoke CUDA and kernels

Some analysis

- ▶ Arguably an order of magnitude faster than the serial method
- ▶ Still some slowdowns in copying to/from GPU, but shouldn't be a problem if clever about it, and not constantly transferring data
- ▶ Slow to invoke CUDA and kernels
- ▶ Timings (from unix time utility, on my laptop):

Some analysis

- ▶ Arguably an order of magnitude faster than the serial method
- ▶ Still some slowdowns in copying to/from GPU, but shouldn't be a problem if clever about it, and not constantly transferring data
- ▶ Slow to invoke CUDA and kernels
- ▶ Timings (from unix time utility, on my laptop):

Some analysis

- ▶ Arguably an order of magnitude faster than the serial method
- ▶ Still some slowdowns in copying to/from GPU, but shouldn't be a problem if clever about it, and not constantly transferring data
- ▶ Slow to invoke CUDA and kernels
- ▶ Timings (from unix time utility, on my laptop):

real 0m5.028s

user 0m0.337s

sys 0m0.087s

Some analysis

- ▶ Arguably an order of magnitude faster than the serial method
- ▶ Still some slowdowns in copying to/from GPU, but shouldn't be a problem if clever about it, and not constantly transferring data
- ▶ Slow to invoke CUDA and kernels
- ▶ Timings (from unix time utility, on my laptop):

real 0m5.028s

user 0m0.337s

sys 0m0.087s

- ▶ note disparity between real and user/sys: very little CPU time, lots of GPU

Some analysis

- ▶ Arguably an order of magnitude faster than the serial method
- ▶ Still some slowdowns in copying to/from GPU, but shouldn't be a problem if clever about it, and not constantly transferring data
- ▶ Slow to invoke CUDA and kernels
- ▶ Timings (from unix time utility, on my laptop):

real 0m5.028s

user 0m0.337s

sys 0m0.087s

- ▶ note disparity between real and user/sys: very little CPU time, lots of GPU
- ▶ however, still faster than serial.

But still not faster than C

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <math.h>
4  int main(int argc, char** argv){
5      float *wav_array = malloc(sizeof(float)*900*900);
6      float pa[] = {300,450};
7      float pb[] = {600,450};
8      int index = 0; float a,b;
9      for(int i = 0;i<900;i++){
10         for(int j = 0;j<900;j++){
11             index = i+(j*900);
12             a = sqrt(((i-pa[0])*(i-pa[0]))+((j-pa[1])*(j-pa[1])));
13             b = sqrt(((i-pb[0])*(i-pb[0]))+((j-pb[1])*(j-pb[1])));
14             wav_array[index] = (sin(a)+sin(b))/2;
15         }
16     }
17     return 0;
18 }
```

C analysis

- ▶ Timings:

C analysis

- ▶ Timings:

C analysis

- ▶ Timings:

real 0m0.110s

user 0m0.107s

sys 0m0.003s

C analysis

- ▶ Timings:

real 0m0.110s

user 0m0.107s

sys 0m0.003s

- ▶ what???

C analysis

- ▶ Timings:

real 0m0.110s

user 0m0.107s

sys 0m0.003s

- ▶ what???
- ▶ not the fault of CUDA, more the fault of python and poor programming

C analysis

- ▶ Timings:

real 0m0.110s

user 0m0.107s

sys 0m0.003s

- ▶ what???
- ▶ not the fault of CUDA, more the fault of python and poor programming
- ▶ for larger problems, with more complex processing, we get a bigger speedup with CUDA

Rewrite the CUDA in C?

```
1  /*Same kernel definition */
2  int main(int argc, char** argv){
3      float *wav_array = (float*)malloc(sizeof(float)*900*900);
4      float *gpu_wav_array;
5      cudaMalloc(&gpu_wav_array, 900*900*sizeof(float));
6      dim3 block_size;
7      block_size.x = 30;
8      block_size.y = 30;
9      block_size.z = 1;
10     dim3 grid_size;
11     grid_size.x = 30;
12     grid_size.y = 30;
13     sin_dist<<<block_size, grid_size>>>(gpu_wav_array);
14     cudaMemcpy(wav_array, gpu_wav_array, 900*900*sizeof(float),
15               cudaMemcpyDeviceToHost);
16     return 0;
17 }
```


Analysis:

- ▶ Timings:

real 0m4.511s

user 0m0.017s

sys 0m0.047s

Analysis:

- ▶ Timings:

real 0m4.511s

user 0m0.017s

sys 0m0.047s

- ▶ Again, slower than C??

Analysis:

- ▶ Timings:

real 0m4.511s

user 0m0.017s

sys 0m0.047s

- ▶ Again, slower than C??
- ▶ yes, what we're doing in the graphics card isn't that hard in serial

Analysis:

- ▶ Timings:

real 0m4.511s

user 0m0.017s

sys 0m0.047s

- ▶ Again, slower than C??
- ▶ yes, what we're doing in the graphics card isn't that hard in serial
- ▶ slowdown switching between CPU and GPU

Analysis:

- ▶ Timings:

real 0m4.511s

user 0m0.017s

sys 0m0.047s

- ▶ Again, slower than C??
- ▶ yes, what we're doing in the graphics card isn't that hard in serial
- ▶ slowdown switching between CPU and GPU
- ▶ however, faster than PyCUDA - so some speedup gaineds

Analysis:

- ▶ Timings:

real 0m4.511s

user 0m0.017s

sys 0m0.047s

- ▶ Again, slower than C??
- ▶ yes, what we're doing in the graphics card isn't that hard in serial
- ▶ slowdown switching between CPU and GPU
- ▶ however, faster than PyCUDA - so some speedup gaineds
- ▶ more importantly, we use the same CUDA code, so if we want to squeeze speed out, we can rewrite in C

Questions!

Any questions?